

SCOTTYRANK.JL: AN IMPLEMENTATION OF PAGERANK & HITS

SIYUAN CHEN AND MICHAEL ZHOU

ABSTRACT. SOME ABSTRACT HERE

1. BACKGROUND

1.1. Linear Algebra.

1.1.1. *Definitions.* Positive matrices are defined as matrices with positive entries.

Markov matrices are defined as square matrices with nonnegative entries and column sum 1 across all of its columns. Note that for a $n \times n$ matrix M , the latter condition is equivalent to $M^T \vec{1} = \vec{1}$, where $\vec{1} \in \mathbb{R}^n$ has all ones as components.

Positive Markov matrices are defined as, well, positive Markov matrices.

1.1.2. *Facts.* (Perron-Frobenius theorem) Let A be a positive square matrix. Let λ_1 be A 's maximum eigenvalue in terms of absolute values. Then λ_1 is positive and has algebraic (and subsequently geometric) multiplicity 1.

Let M be a Markov matrix. Let λ_1 be M 's maximum eigenvalue in terms of absolute values. Then $\lambda_1 = 1$.

Let M' be a positive Markov matrix. Let λ_1 be M' 's maximum eigenvalue in terms of absolute values. Then $\lambda_1 = 1$ and has algebraic (and subsequently geometric) multiplicity 1.

1.1.3. *Usage.* Let M be a $n \times n$ Markov matrix. Then M specifies a discrete memoryless transition process between n states, namely the process where

$$(\forall (t, i, j) \in \mathbb{N} \times [n] \times [n]) [\Pr(\text{state } i \text{ at time } t+1 \mid \text{state } j \text{ at time } t) = M_{ij}].$$

Let $\vec{v} \in \mathbb{R}^n$ such that \vec{v} has nonnegative components and $\vec{v}^T \vec{1} = 1$ (a stochastic vector). Then \vec{v} specifies an (initial) discrete probability distribution over the n states, namely the distribution where

$$(\forall i \in [n]) [\Pr(\text{state } i \text{ at time } 0) = \vec{v}_i].$$

Then the probability distribution over the n states after t steps of the transition process specified by M is precisely $M^t \vec{v}$, or equivalently

$$(\forall (t, i) \in \mathbb{N} \times [n]) [\Pr(\text{state } i \text{ at time } t) = (M^t \vec{v})_i].$$

1.2. Graph Theory.

1.2.1. *Definitions.* A simple directed graph is defined as an unweighted directed graph without self-referential edges or multiple edges between the same origin destination pair.

For a simple directed graph with n vertices, the adjacency matrix \mathcal{A} is defined to be the $n \times n$ matrix where

$$(\forall (i, j) \in [n] \times [n]) \left(A_{ij} = \begin{cases} 1 & \text{there is an edge to } i \text{ from } j \\ 0 & \text{otherwise} \end{cases} \right).$$

1.2.2. *Facts.* For a simple directed graph with n vertices and its adjacency matrix \mathcal{A} ,

$$\begin{aligned} (\forall j \in [n]) & \left[\text{number of outgoing neighbors from vertex } j = \text{out}(j) = (\mathcal{A}_{*j})^T \vec{1} \right] \\ (\forall i \in [n]) & \left[\text{number of incoming neighbors to vertex } i = \text{in}(i) = (\mathcal{A}_{i*})^T \vec{1} \right]. \end{aligned}$$

2. ALGORITHMS

2.1. The Network Model. Both algorithms, PageRank and HITS, model the network of interest as a simple directed graph with websites as vertices and links as edges. This implies that there will be no self-referential links, no duplicate links between the same origin and destination pair, and no priority difference between links.

2.2. PageRank.

2.2.1. The random walk. PageRank models the behavior of a typical web surfer as a damped random walk.

- (1) The surfer starts out by visiting a random site out of all sites with equal probability.
- (2) At every step, the surfer has a probability λ of continuing surfing and a complementary $1 - \lambda$ probability of losing interest, for a predetermined λ .
 - (a) If the surfer continues ...
 - (i) ... and there are links exiting the current site, the surfer clicks on a random link (and visits the site it points to) out of those links with equal probability.
 - (ii) ... and there aren't any links exiting the current site, the surfer simply visits a random site out of all other sites with equal probability.
 - (b) If the surfer loses interest, they simply visits a random site out of all sites with equal probability.

To best model a typical surfer's probability of continuing surfing, λ , also known as the damping factor, is empirically determined to be around 0.85.

2.2.2. Matrix representation. Let n be the number of websites in the network of interest. Let \mathcal{A} be the adjacency matrix for the network of interest. Let $\langle \vec{v}_t \rangle_{t \in \mathbb{N}}$ be the probability distributions describing the website the surfer is visiting at time t . Let M be the transition matrix for the random walk process.

Then $\vec{v}_0 = \vec{1}/n$, M is the $n \times n$ matrix where

$$(\forall (i, j) \in [n] \times [n]) \left[M_{ij} = \begin{cases} \frac{\lambda}{\text{out}(j)} + \frac{1-\lambda}{n} & \mathcal{A}_{ij} = 1 \\ \frac{1-\lambda}{n} & \mathcal{A}_{ij} = 0 \wedge \text{out}(j) > 0 \\ \frac{\lambda}{n-1} + \frac{1-\lambda}{n} & i \neq j \wedge \text{out}(j) = 0 \\ \frac{1-\lambda}{n} & i = j \wedge \text{out}(j) = 0 \end{cases} \right],$$

and

$$(\forall t \in \mathbb{N}) (\vec{v}_t = M^t \vec{v}_0).$$

Note that in this case M is a positive Markov matrix, assuming reasonable λ .

2.2.3. Definition. The PageRank score for a given website in the network of interest is defined as the probability of a typical surfer visiting that website after an indefinitely long damped random walk. In matrix form,

$$(\forall i \in [n]) \left[\text{PageRank}(i) = \lim_{t \rightarrow \infty} (\vec{v}_t)_i = \lim_{t \rightarrow \infty} (M^t \vec{v}_0)_i \right].$$

Note that the limits exist: convergence is guaranteed as M has a unique maximal eigenvalue of 1 and thus an steady attracting state.

2.3. HITS.

2.3.1. Authorities and hubs. Due to PageRank’s algorithmic design, a given website’s PageRank score determined mostly by the scores of its incoming neighbors. Consequently, PageRank tends to underestimate the importance of websites similar to “web directories”, i.e., those with few significant incoming neighbors yet many significant outgoing neighbors.

To address this issue, HITS (Hyperlink-Induced Topic Search) introduces Authority and Hub scores, which measure a given website’s tendencies to be referred to by others and to refer to others, respectively. Note that the two metrics are not “mutually exclusive”; a website like Wikipedia can have both a high Authority score and a high Hub score.

Specifically, Authority and Hub scores are recursively defined: a website’s Authority score is determined by the Hub scores of its incoming neighbors and its Hub score is determined by the Authority scores of its outgoing neighbors.

2.3.2. Matrix representation. Let n be the number of websites in the network of interest. Let \mathcal{A} be the adjacency matrix for the network of interest. Let $\langle \vec{a}_t \rangle_{t \in \mathbb{N}}$ and $\langle \vec{h}_t \rangle_{t \in \mathbb{N}}$ be the (pre-normalization) Authority and Hub scores for the n websites at time t .

Then $\vec{a}_0 = \vec{h}_0 = \vec{1}$ and

$$(\forall t \in \mathbb{N}) \left[\left(\vec{a}_{t+1}, \vec{h}_{t+1} \right) = \left(\mathcal{A} \vec{h}_t, \mathcal{A}^T \vec{a}_t \right) \right].$$

2.3.3. Definition. The Authority and Hub scores for a given website in the network of interest is defined as the respective scores after indefinitely many iterations. In matrix form,

$$(\forall i \in [n]) \left[(\text{Authority}(i), \text{Hub}(i)) = \lim_{t \rightarrow \infty} \left((\vec{a}_t)_i, (\vec{h}_t)_i \right) \right].$$

To guarantee convergence, the Authority and Hub scores are normalized. Our implementation performs normalization after every iteration. This means

$$(\forall t \in \mathbb{N}) \left[\|(\vec{a}_t)'\| = \|(\vec{h}_t)'\| = 1 \right]$$

where

$$(\forall t \in \mathbb{N}) \left[(\vec{a}_t)' = \frac{\vec{a}_t}{\|\vec{a}_t\|} \wedge (\vec{h}_t)' = \frac{\vec{h}_t}{\|\vec{h}_t\|} \right].$$

3. IMPLEMENTATION

3.1. Structs. We define two structs, **Vertex** and **Graph**, to represent the vertices and the graph itself in our simple directed graph model for the network of interest.

Note that to align with Julia conventions, we use 1-based indexing.

```

1  # export Vertex, Graph
2
3  struct Vertex
4      index::UInt32
5      in_neighbors::Vector{UInt32}
6      out_neighbors::Vector{UInt32}

```

```

7  end
8
9  struct Graph
10     num_vertices::UInt32
11     vertices::Vector{Vertex}
12 end

```

3.2. Input. We define three functions, `read_graph`, `read_edge_list`, and `read_adjacency_list`, to read and construct graphs from text files. We expose `read_graph` to the client with the option to specify the type of the input file and whether or not the input file uses 0-based indexing.

Edge list files follow the following format:

```

1  [num_nodes] [num_edges]
2  [index_from] [index_to] # repeats [num_edges] times
3  ...                  # in total

```

Adjacency list files follow the following format:

```

1  [num_nodes]
2  [index_to_1] [index_to_2] ... [index_to_m] # repeats [num_nodes] times
3  ...                  # in total

```

The code for `read_graph`, `read_edge_list`, and `read_adjacency_list` can be found in the Appendix.

3.3. PageRank. We divide the PageRank algorithm into three steps:

- (1) Generating the transition matrix: `pagerank_matrix`.
- (2) Running the transition process: `pagerank_iteration`, `pagerank_epsilon`.
- (3) Returning the desired output: `pagerank_print`, `pagerank`.

3.3.1. Generating the transition matrix. The function `pagerank_matrix` generates a Markov matrix M that specifies the transition probabilities of the PageRank transition process.

We first compute the entries in M prior to damping, casing on whether the origin vertex is a “sink” (no outgoing neighbors), and then apply the damping at the end.

```

1  function pagerank_matrix(graph::Graph, damping::Float64)
2      M = zeros(Float64, (graph.num_vertices, graph.num_vertices))
3      for vertex in graph.vertices
4          num_out_neighbors = length(vertex.out_neighbors)
5          if num_out_neighbors == 0
6              for index_to in 1:graph.num_vertices
7                  M[index_to, vertex.index] = 1 / (graph.num_vertices - 1)
8              end
9              M[vertex.index, vertex.index] = 0
10         else
11             for index_to in vertex.out_neighbors
12                 M[index_to, vertex.index] = 1 / num_out_neighbors
13             end
14         end

```

```

15     end
16     map(x -> damping * x + (1 - damping) / graph.num_vertices, M)
17 end

```

3.3.2. *Running the transition process.* The functions `pagerank_iteration` and `pagerank_epsilon` both generate an initial stochastic vector and then carry out the transition process using the transition matrix.

`pagerank_iteration` runs the process for a given number of iterations.

```

1 function pagerank_iteration(num_vertices::UInt32, M::Matrix{Float64},
   ↪ num_iterations::UInt32)
2     M_pwr = Base.power_by_squaring(M, num_iterations)
3     M_pwr * (ones(Float64, num_vertices) / num_vertices)
4 end

```

`pagerank_epsilon` runs the process until the norm of the difference vector is smaller than a given threshold, or until $\|\vec{v}_{k+1} - \vec{k}\| < \epsilon$.

```

1 function pagerank_epsilon(num_vertices::UInt32, M::Matrix{Float64},
   ↪ epsilon::Float64)
2     prev = ones(Float64, num_vertices) / num_vertices
3     curr = M * prev
4     while norm(prev - curr) > epsilon
5         prev, curr = curr, M * curr
6     end
7     curr
8 end

```

3.3.3. *Returning the desired output.* We expose two functions to the client: `pagerank` and `pagerank_print`.

`pagerank` calculates the PageRank scores for the input `graph`, with the option to specify the damping factor and the transition mode.

```

1 # export pagerank_print, pagerank
2
3 function pagerank(graph::Graph;
4     damping::Float64=0.85, modeparam::Tuple{String, Union{Int64, UInt32,
   ↪ Float64}}=("iter", 10))
5     if damping < 0 || damping > 1
6         error("invalid damping")
7     end
8     M = pagerank_matrix(graph, damping)
9     if modeparam[1] == "iter"
10         if !(isinteger(modeparam[2])) || modeparam[2] < 0
11             error("invalid param")
12         end
13         pagerank_iteration(graph.num_vertices, M, UInt32(modeparam[2]))
14     elseif modeparam[1] == "epsi"
15         if modeparam[2] <= 0
16             error("invalid param")
17         end
18     end
19 end

```

```

18     pagerank_epsilon(graph.num_vertices, M, Float64(modeparam[2]))
19 else
20     error("invalid mode")
21 end
22 end

```

`pagerank_print` pretty-prints the PageRank scores along with relevant information about the top vertices for the input `graph` and scores `pg`.

The code for `pagerank_print` can be found in the Appendix.

3.4. **HITS.** Similarly, we divide the HITS algorithm into three steps:

- (1) Generating the transition matrix: `hits_matrix`.
- (2) Running the transition process: `hits_update`, `hits_iteration`, `hits_epsilon`.
- (3) Returning the desired output: `hits_print`, `hits`.

3.4.1. *Generating the transition matrix.* The transition matrices for the hits algorithm are simply the adjacency matrix and its transpose.

```

1 function hits_matrix(graph::Graph)
2     A = zeros(Float64, (graph.num_vertices, graph.num_vertices))
3     for vertex in graph.vertices
4         for index_to in vertex.out_neighbors
5             A[index_to, vertex.index] = 1
6         end
7     end
8     A
9 end

```

3.4.2. *Running the transition process.* The function `hits_update` computes the normalized new Authority and Hub scores from the previous Authority and Hub scores and the two transition matrices

```

1 function hits_update(A::Matrix{Float64}, H::Matrix{Float64},
2     ↪ a::Vector{Float64}, h::Vector{Float64})
3     normalize(A * h), normalize(H * a)
4 end

```

The functions `hits_iteration` and `hits_epsilon` both generate the initial Authority and Hub scores and then carry out the transition process using the update function.

`hits_iteration` runs the process for a given number of iterations.

```

1 function hits_iteration(num_vertices::UInt32, A::Matrix{Float64},
2     ↪ H::Matrix{Float64}, num_iterations::UInt32)
3     a, h = ones(Float64, num_vertices), ones(Float64, num_vertices)
4     for _ in 1:num_iterations
5         a, h = hits_update(A, H, a, h)
6     end
7     a, h
8 end

```

`hits_epsilon` runs the process until the norms of both difference vectors are smaller than the given threshold, or until $\|\vec{a}_{k+1} - \vec{a}_k\| < \epsilon \wedge \|\vec{h}_{k+1} - \vec{h}_k\| < \epsilon$.

```

1 function hits_epsilon(num_vertices::UInt32, A::Matrix{Float64},
  ↪ H::Matrix{Float64}, epsilon::Float64)
2   prev_a, prev_h = ones(Float64, num_vertices), ones(Float64,
  ↪ num_vertices)
3   curr_a, curr_h = hits_update(A, H, prev_a, prev_h)
4   while norm(prev_a - curr_a) > epsilon || norm(prev_h - curr_h) >
  ↪ epsilon
5     prev_a, prev_h, (curr_a, curr_h) = curr_a, curr_h, hits_update(A, H,
  ↪ curr_a, curr_h)
6   end
7   curr_a, curr_h
8 end

```

3.4.3. *Returning the desired output.* We expose two functions to the client: `hits` and `hits_print`

`hits` calculates the Authority and Hub scores for the input `graph`, with the option to specify the transition mode.

```

1 # export hits_print, hits
2
3 function hits(graph::Graph;
4   modeparam::Tuple{String, Union{Int64, UInt32, Float64}}=("iter", 10))
5   A = hits_matrix(graph)
6   H = copy(transpose(A))
7   if modeparam[1] == "iter"
8     if !(isinteger(modeparam[2])) || modeparam[2] < 0
9       error("invalid param")
10    end
11    hits_iteration(graph.num_vertices, A, H, UInt32(modeparam[2]))
12  elseif modeparam[1] == "epsi"
13    if modeparam[2] <= 0
14      error("invalid param")
15    end
16    hits_epsilon(graph.num_vertices, A, H, Float64(modeparam[2]))
17  else
18    error("invalid mode")
19  end
20 end

```

`hits_print` pretty-prints the Authority and Hub scores along with relevant information about the top vertices for the input `graph` and scores `a` and `h`.

The code for `hits_print` can be found in the Appendix.

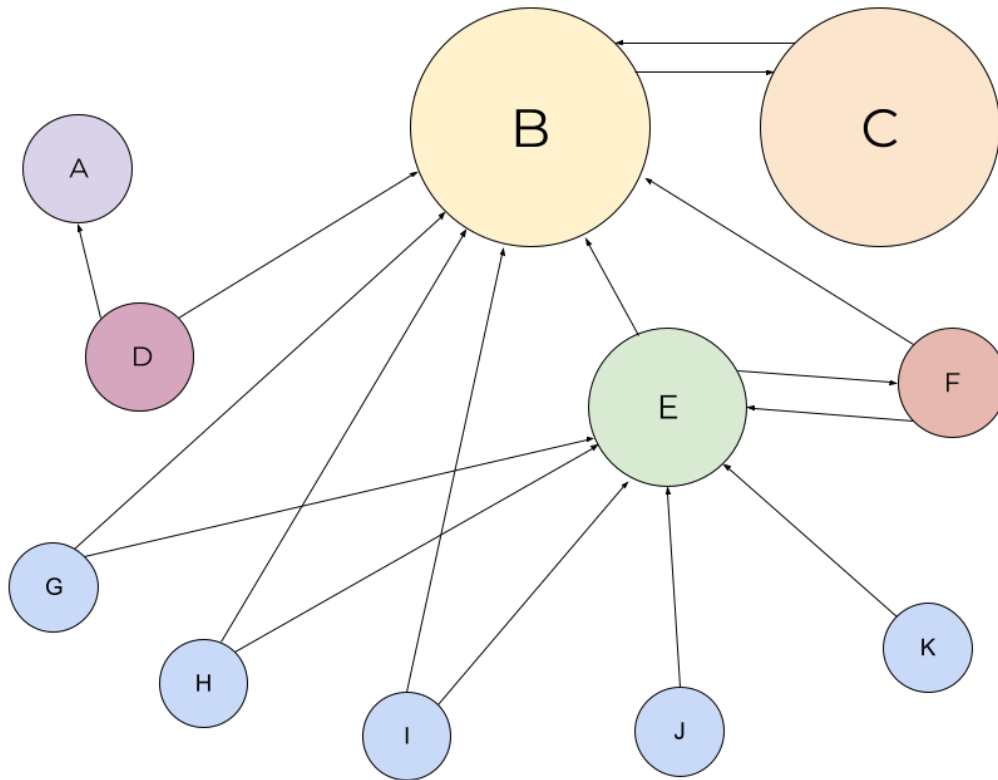
3.5. **Output.** We offer two ways of exporting the `graph` struct: `generate_adjacency_matrix` and `generate_adjacency_list`.

`generate_adjacency_matrix`, well, generates the `graph`'s adjacency matrix, with the option to specify whether the desired output should use 0-based indexing.

The code for `generate_adjacency_matrix` and `generate_adjacency_list` can be found in the Appendix.

4. RESULTS

Shown below is a recreation of the medium-sized network with 11 vertices displayed in the Wikipedia article on PageRank:



We will run both PageRank and HITS on the network above.

4.1. PageRank. We expect vertices B and C to have the highest PageRank scores because virtually every vertex points to it, either directly or through other vertices such as E. We also expect E to have a fairly high PageRank score because there are many nodes pointing to it; however, E points to B but B does not point to E, so E should have a lower PageRank than B. Furthermore, we expect the vertices G, H, I, J, and K to have the lowest PageRank scores because no vertices point towards them, and they only have edges pointing outwards. Their PageRank scores should also be roughly equal since they each have a similar outgoing edge structure, only pointing to 1 or 2 other vertices.

We run the following commands to apply PageRank on the graph.

```

1 $ julia
2 julia> using ScottyRank
3 julia> G=read_graph("data/medium-el.txt", "el") # the graph is stored as
  ↳ an edge list
4 julia> pg = pagerank(G) # run PageRank on the graph
5 julia> pagerank_print(G, pg, num_lines=11)

```

The printed result for `pagerank_print` is shown below:

PageRank Score	Index	Vertex	In	Out
0.3643	2	B	7	1
0.3638	3	C	1	1
0.0813	5	E	6	3
0.0395	4	D	1	2
0.0395	6	F	1	2
0.0304	1	A	1	0
0.0163	9	I	0	2
0.0163	10	J	0	1
0.0163	11	K	0	1
0.0163	7	G	0	2
0.0163	8	H	0	2

The first column lists the calculated PageRank score and the second lists the index of the vertex with that score. The third, titled “Vertex”, is the letter assigned to that vertex (A for index 1, B for index 2, etc.). The “In” column displays the number of inbound edges to the corresponding vertex. Similarly, the “Out” columns displays the number of outbound edges from the corresponding vertex.

We see that, as expected, vertices B and C have the highest PageRank scores by far, and E has the highest score of the remaining vertices. Additionally, we see the vertices G, H, I, and J share the same lowest PageRank score of 0.0163.

4.2. HITS on Medium-Sized Graph. When predicting the results for HITS on the graph, we must recall the concepts behind HITS, namely that a vertex pointing to vertices with high authority scores will have a high hub score, and that a vertex which has vertices with high hub scores pointing to it will have a high authority score.

With this in mind, we see that vertices B and E should have the highest authority scores simply because they have the greatest amount of vertices pointing towards them.

However, since B does not point towards many other authorities, its hub score should be low. But E points to B, which is a high authority score vertex, so we should expect the hub score of E to be above average.

Similarly, the hub score of C should be relatively high since it points to B, but the authority score of C should be low because B does not have a high hub score.

Finally, we should expect the authority scores of G, H, I, J, and K all to be very low because no vertices point towards them, but their hub scores to be relatively high because they point to one or both of the highest authority nodes– B and E.

We run the following commands to apply HITS on the graph.

```

1 $ julia
2 julia> using ScottyRank
3 julia> G=read_graph("data/medium-el.txt", "el") # the graph is stored as
  ↪ an edge list
4 julia> a,h = hits(G) # run HITS on the graph
5 julia> hits_print(G, a, h, num_lines=11)

```

The printed result for `hits_print` is shown below:

Authority Score	Index	Vertex	In	Out
0.7554	2	B	7	1
0.6388	5	E	6	3
0.0870	4	D	1	2
0.0870	6	F	1	2
0.0779	1	A	1	0
0.0000	3	C	1	1
0.0000	7	G	0	2
0.0000	8	H	0	2
0.0000	9	I	0	2
0.0000	10	J	0	1
0.0000	11	K	0	1

Hub Score	Index	Vertex	In	Out
0.4259	6	F	1	2
0.4259	7	G	0	2
0.4259	8	H	0	2
0.4259	9	I	0	2
0.2835	5	E	6	3
0.2543	4	D	1	2
0.2306	3	C	1	1
0.1953	10	J	0	1
0.1953	11	K	0	1
0.0000	2	B	7	1
0.0000	1	A	1	0

Like PageRank, the first column lists the calculated authority score for the first table and the hub score for the second table. The second column lists the index of the vertex with that score. The third, titled “Vertex”, is the letter assigned to that vertex (A for index 1, B for index 2, etc.). The “In” column displays the number of inbound edges to the corresponding vertex. Similarly, the “Out” columns displays the number of outbound edges from the corresponding vertex.

As predicted, the authority scores of B and E are the highest, with C having a low authority score. The vertices G, H, I, J, and K have authority scores of 0 because no vertices point to them.

Additionally, we see that the hub scores were also calculated as predicted— the hub score of F was the highest because F pointed towards both B and E, the two vertices with the highest authority scores. As justified in our predictions, G, H, I have very high hub scores. However, J and K had lower hub scores because they only pointed to one of the main authorities.